

Haizhou Shi

Updated November 14, 2025

Email: haizhou.shi.057@gmail.com

Homepage: haizhou-shi.github.io

Phone: (+1) 7324334667

Research Interests	Bayesian LLMs and Continual Training of LLMs.	
Education	Rutgers University Ph.D in Computer Science, focusing on Machine Learning Advisor: Professor Hao Wang .	New Jersey, US Sept. 2022 – Present
	Zhejiang University B.Eng & M.Eng in Computer Science and Technology Mentors: Professors Siliang Tang , Yueling Zhuang .	Hangzhou, China Sept. 2015 – June 2022
Under Review	(2025) SSR: Socratic Self-Refine for Large Language Model Reasoning. Haizhou Shi , Ye Liu, Bo Pang, Zeyu Leo Liu, Hao Wang, Silvio Savarese, Caiming Xiong, Yingbo Zhou, Semih Yavuz. arXiv preprint at arXiv:2511.10621	
	(2025) Token-Level Uncertainty Estimation for Large Language Model Reasoning. Tunyu Zhang*, Haizhou Shi *, Yibin Wang, Hengyi Wang, Xiaoxiao He, Zhuowei Li, Haoxian Chen, Ligong Han, Kai Xu, Huan Zhang, Dimitris Metaxas, Hao Wang. NeurIPS-W'25, arXiv preprint at arXiv:2505.11737	
	(2025) Efficient Uncertainty Estimation via Distillation of Bayesian LLMs. Harshil Vejendla*, Haizhou Shi *, Yibin Wang, Tunyu Zhang, Huan Zhang, Hao Wang. arXiv preprint at arXiv:2505.11731	
Publication	(2025) Training-Free Bayesianization for Low-Rank Adapters of Large Language Models. Haizhou Shi *, Yibin Wang*, Ligong Han, Huan Zhang, Hao Wang. NeurIPS'25, arXiv preprint at arXiv:2412.05723	
	(2025) The Hidden Life of Tokens: Reducing Hallucination of Large Vision-Language Models via Visual Information Steering. Zhuowei Li, Haizhou Shi , Yunhe Gao, Di Liu, Zhenting Wang, Yuxiao Chen, Ting Liu, Long Zhao, Hao Wang, Dimitris N. Metaxas. ICML'25, arXiv preprint at arXiv:2502.03628	
	(2025) Continual Learning of Large Language Models: A Comprehensive Survey. Haizhou Shi , Zihao Xu, Hengyi Wang, Weiyi Qin, Wenyuan Wang, Yibin Wang, Zifeng Wang, Sayna Ebrahimi, Hao Wang. ACM Computing Surveys, arXiv preprint at arXiv:2404.16789	
	(2024) BLoB: Bayesian Low-Rank Adaptation by Backpropagation for LLMs. Yibin Wang*, Haizhou Shi *, Ligong Han, Dimitris Metaxas, Hao Wang. NeurIPS'24, arXiv preprint at arXiv:2406.11675	
	(2023) A Unified Approach to Domain Incremental Learning with Memory: Theory and Algorithm. Haizhou Shi , Hao Wang. NeurIPS'23, arXiv preprint arXiv:2310.12244	

(2021) **On the Efficacy of Small Self-Supervised Contrastive Models without Distillation Signals.**

Haizhou Shi, Youcai Zhang, Siliang Tang, Wenjie Zhu, Yaqian Li, Yandong Guo, Yueling Zhuang.

AAAI'22 oral, arXiv preprint arXiv:2107.14762

(2021) **Towards Communication-Efficient and Privacy-Preserving Federated Representation Learning.**

Haizhou Shi, Youcai Zhang, Zijin Shen, Siliang Tang, Yaqian Li, Yandong Guo, Yueling Zhuang.

FL-AAAI-22 oral, arXiv preprint arXiv:2109.14611

(2021) **CIL: Contrastive Instance Learning Framework for Distantly Supervised Relation Extraction.**

Tao Chen, **Haizhou Shi**, Siliang Tang, Zhigang Chen, Fei Wu, and Yueling Zhuang.

ACL '21, arXiv preprint arXiv:2106.10855

(2019) **Informative Visual Storytelling with Cross-Modal Rules.**

Jiacheng Li, **Haizhou Shi**, Siliang Tang, Fei Wu, Yueling Zhuang.

ACM-MM'19, Proceedings of the 27th ACM International Conference on Multimedia

Work Experience

AI Research Intern, at Salesforce AI Research.

Palo Alto, US

Mentors: [Semih Yavuz](#), [Yingbo Zhou](#).

June 2025 – Present

Led the research project on LLM reasoning.

Led the research project on Reward Modeling of LLMs.

Machine Learning Research Intern, at Morgan Stanley.

New York, US

Mentors: [Marin Bilos](#), [Kashif Rasul](#).

June 2024 – Aug. 2024

Led the research project on lifelong time series forecasting.

Research Intern, at OPPO Research Institute.

Shanghai, China

Mentors: [Youcai Zhang](#), [Yandong Guo](#).

Dec. 2020 – Nov. 2021

Led the research on self-supervised learning for lightweight backbones (paper available [here](#)).

Led the research on federated self-supervised learning (paper available [here](#)).

Research Intern, at Intel Asia-Pacific Research & Development LTD.

Shanghai, China

Mentors: [Shengsheng Huang](#), [Jinquan Dai](#).

Apr. 2019 – Aug. 2019

Conducted reinforcement learning research on the google football simulated environment, benchmarked multiple existing frameworks including DQN, PG, and IMPALA.

Service & Outreach

Life-long Fellow, at Melton Foundation.

June 2016 – Present

Learned and practiced the concept of the global citizenship. Participated multiple global conferences on promoting and enabling global citizenship as a way for individuals and organizations to work together across boundaries of place and identity to address global challenges.

Reviewer, at Academic Conferences

Provided review services for multiple top-tier ML&AI conferences, including *ICLR'26, NeurIPS'25 (Top-Reviewer Award), ICML'25, ICLR'25, ACL ARR'25, AAAI'25, NeurIPS'24, ECCV'24, IJCAI'24, AAAI'24, AAAI'23, ICCV'23, ECCV'22*.